

Neural Solvers for Fast and Accurate Numerical Optimal Control

Federico Berto¹, Stefano Massaroli²
Michael Poli³, Jinkyoo Park¹

¹KAIST, ²The University of Tokyo, ³Stanford University

*35th Conference on Neural Information Processing Systems
NeurIPS 2021 DLDE Workshop*



Introduction

Optimal Control Problem: find controller parameters θ to minimize cost

$$\boxed{\begin{array}{ll} \min_{\theta} & \mathbb{E}_{x_0 \sim \rho_0(x_0)} [\mathcal{J}(x_0, u_{\theta}(t))] \\ \text{subject to} & \dot{x}(t) = f(t, x(t), u_{\theta}(t)) \\ & x(0) = x_0; \quad t \in \mathcal{T} \end{array}}$$

with cost functional

$$\mathcal{J} = [x^T(t_f) - x^{\star}] \mathbf{P} [x(t_f) - x^{\star}] + \int_{t_0}^{t_f} \left([x^T(t) - x^{\star}] \mathbf{Q} [x(t) - x^{\star}] + u_{\theta}^T(t) \mathbf{R} u_{\theta}(t) \right) dt$$

Key considerations:

- **Accuracy** is essential to obtain correct solutions
- **Speed** is crucial for hard-time constrained applications
- Fast but inaccurate solvers **trade off** solution accuracy for speed

Introduction

Explicit ODE Solvers: iterate to solve

$$x_{k+1} = x_k + \epsilon \psi_\epsilon(t_k, x_k, u_k)$$

- ODE solvers differ in how the map ψ is designed
- *Higher-order* (explicit) solvers compute ψ iteratively in p steps
(p denotes the order of the solver)

Solver Residuals: given nominal solutions Φ of the optimal control problem

$$R_k = \frac{1}{\epsilon^{p+1}} \left[\Phi(x_k, t_k, t_{k+1}) - x_k - \epsilon \psi_\epsilon(t_k, x_k, u_k) \right]$$

Hypersolvers for Optimal Control

We extend the paradigm of *hypersolvers*¹ to controlled systems

Hypersolvers for Optimal Control

Control input $\textcolor{teal}{u}$ becomes a key component in hypersolver design

$$x_{k+1} = \underbrace{x_k + \epsilon \psi_\epsilon(t_k, x_k, \textcolor{teal}{u}_k)}_{\text{base solver step}} + \epsilon^{p+1} \underbrace{g_\omega(t_k, x_k, \textcolor{teal}{u}_k)}_{\text{hypersolver}}$$

Residual Fitting: training of g_ω can be performed by minimizing

$$\ell_\omega = \frac{1}{K} \sum_{k=0}^{K-1} \|R_k - g_\omega(t_k, x_k, \textcolor{teal}{u}_k)\|_2$$

⇒ Hypersolvers have **theoretical guarantees**¹ on error bounds

¹Poli et al., *Hypersolvers: Toward Fast Continuous-Depth Models*, NeurIPS 2020

Hypersolver Training

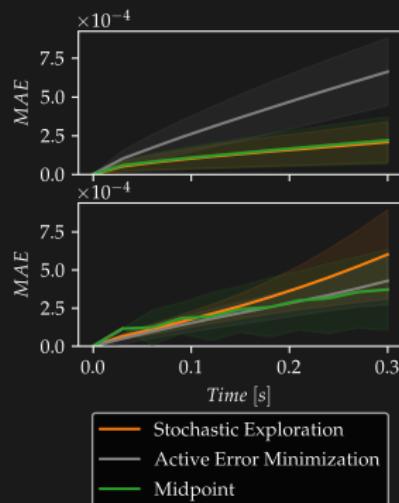
Exploration Strategies: we train hypersolvers on a distribution ξ of states and controllers

- Directly optimize the hypersolver by sampling from the distribution (*stochastic exploration*)

$$\min_w \mathbb{E}_{(x,u) \sim \xi(x,u)} \ell_w(t, x, u)$$

- We can also prioritize exploration in regions with higher residuals (*active error minimization*)

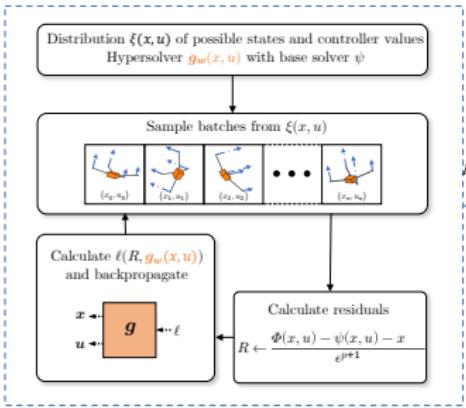
$$\min_w \max_{u \in \mathcal{U}} \mathbb{E}_{(x,u) \sim \xi(x,u)} \ell_w(t, x, u)$$



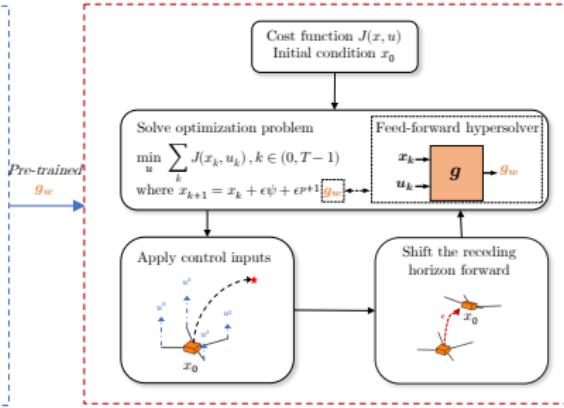
Blueprint of our Approach

- ① **Hypersolver Training:** *offline* optimization of neural approximator
- ② **Optimal Control:** *online* optimization of the controller with increased accuracy and hypersolver low overhead

Hypersolver Training

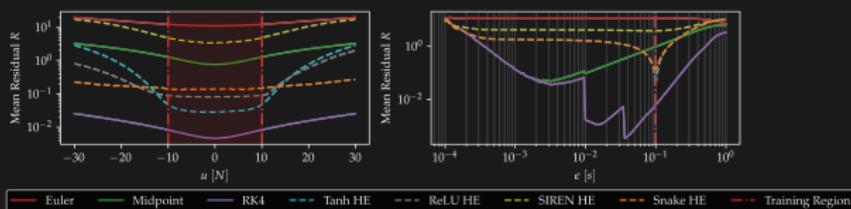
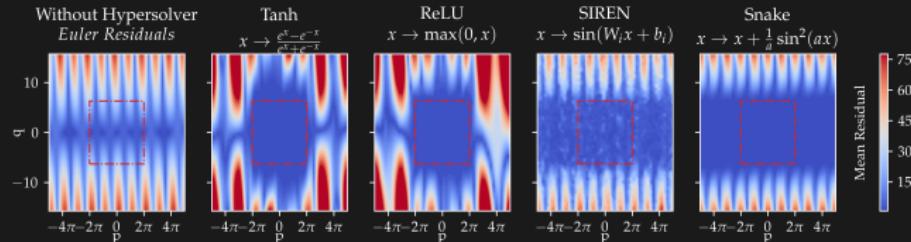


Optimal Control



Architectures

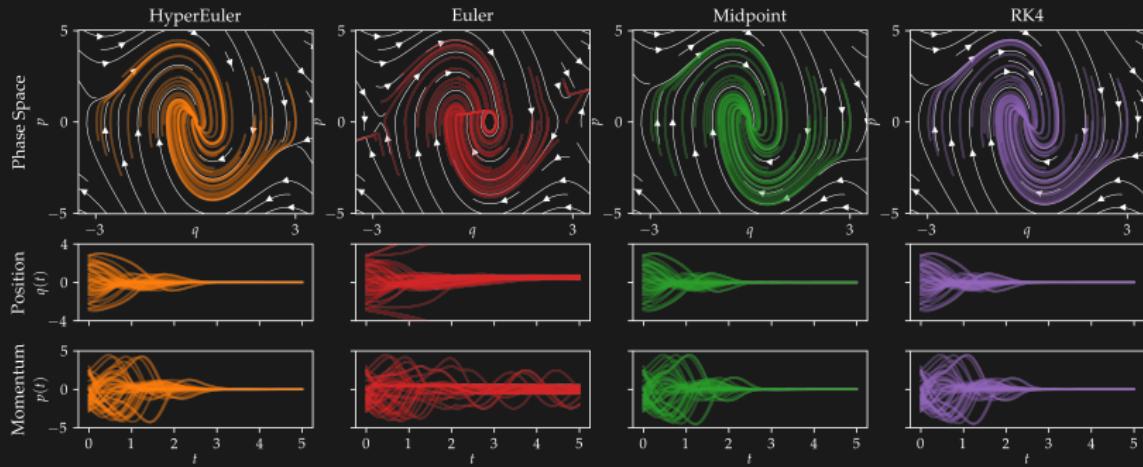
We study different **activation functions** as a core hypersolver design choice



- Different hypersolvers fit the state training region while the ones with *sinusoidal components* also fit better than the base solver outside of it
- Hypersolvers perform better than base solvers even for unseen controllers and timesteps

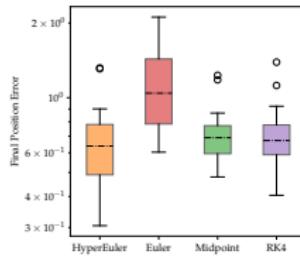
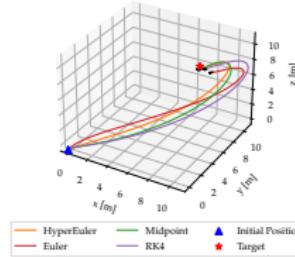
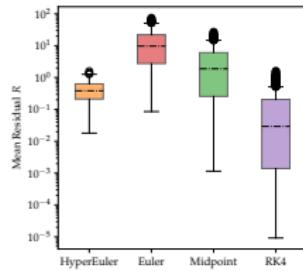
Pendulum Control

Pendulum Controlled Trajectories



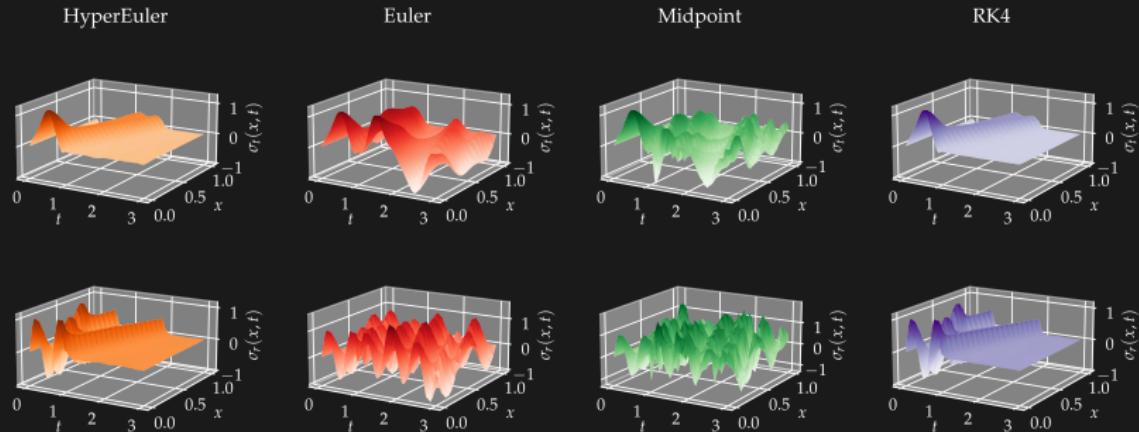
- Task: stabilize the pendulum in the vertical position
- Policy comparable to higher-order solvers with almost half the FLOPs

Quadcopter Model Predictive Control



- Task: move quadcopter to a target position
- Lower residual and position error than Midpoint with half the NFE

Timoshenko PDE Boundary Control



- **Task:** stabilize the beam in resting position
- We control the beam with **HyperEuler** with 70% less runtime w.r.t. **RK4**
- Hypersolvers are even more impactful in high-dimensional settings

Hypersolvers for Optimal Control

Neural solvers for numerical optimal control

- Accuracy improvement with reduced computational cost
- Better downstream control policy
- Even more crucial for high-dimensional control

Some Open Questions:

- How well do hypersolvers perform outside of simulation?
- Can we generalize to different systems?
- What is the most efficient *implicit representation* of hypersolver models?

Thank You

Do not hesitate to contact us for further clarifications!



[DiffEqML](#): Open research community working at the intersection between machine learning, differential equations and control. Feel free to join our Slack channel!