

Learning Optimal Control Policies with Ordinary Differential Equations

Federico Berto

Lab Seminar



SYSTEM INTELLIGENCE LAB



Background: Continuous-depth Framework

Continuous-depth (time) framework

Analyze and develop models bridging the gap between
dynamical system theory and **deep learning**

deep learning \Rightarrow dynamical systems
dynamical systems \Rightarrow deep learning

Differential equations are everywhere: **physics, engineering, sciences.**

A Review on Neural ODEs

Neural ODE: a Core Primitive

We seek the **deep limit** of neural networks

⇒ The input–output map is realized by the *flow* of an ODE¹

Neural ODE [Sonoda, et al. 2017, Chen et al., 2018]

By noticing that the **latent dynamics** of a ResNet

$$\mathbf{z}_{s+1} = \mathbf{z}_s + f_{\theta_s}(\mathbf{z}_s)$$

resemble the *Euler* discretization

$$\frac{\mathbf{z}_{s+1} - \mathbf{z}_s}{\Delta s} \approx \frac{d\mathbf{z}}{ds} = f_{\theta_s}(\mathbf{z}_s) \quad (\Delta s = 1)$$

of the ODE

$$\begin{aligned} \frac{d\mathbf{z}}{ds} &= f_{\theta_s}(s, \mathbf{z}(s)) \\ \mathbf{z}(0) &= \mathbf{x} \end{aligned} \quad s \in \mathcal{S} \subset \mathbb{R}$$

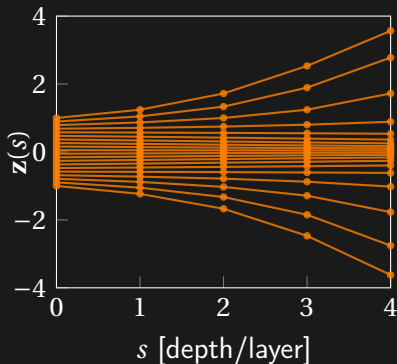
¹Ordinary Differential Equation

Neural ODEs vs ResNets


Inference Model

$$\hat{\mathbf{y}} = \mathbf{x} + \int_0^S f_{\theta}(s, \mathbf{z}(s)) ds$$

ResNet



Neural ODE



A General Neural ODE Formulation

Neural Ordinary Differential Equation

$$\begin{cases} \dot{\mathbf{z}}(s) = f_{\theta(s)}(s, \mathbf{x}, \mathbf{z}(s)) \\ \mathbf{z}(0) = h_x(\mathbf{x}) \\ \hat{\mathbf{y}}(s) = h_y(\mathbf{z}(s)) \end{cases} \quad s \in \mathcal{S}$$

Input	\mathbf{x}	\mathbb{R}^{n_x}
Output	$\hat{\mathbf{y}}$	\mathbb{R}^{n_y}
(Hidden) State	\mathbf{z}	\mathbb{R}^{n_z}
Parameters	$\theta(s)$	\mathbb{R}^{n_θ}
Neural Vector Field	$f_{\theta(s)}$	\mathbb{R}^{n_z}
Input Network	h_x	$\mathbb{R}^{n_x} \rightarrow \mathbb{R}^{n_z}$
Output Network	h_y	$\mathbb{R}^{n_z} \rightarrow \mathbb{R}^{n_y}$

where f_θ , h_x , h_y are generally neural networks.

Inference of general Neural ODEs

$$\hat{\mathbf{y}}(S) = h_y \left(h_x(\mathbf{x}) + \int_0^{S_x^*} f_{\theta(s)}(s, \mathbf{x}, \mathbf{z}(s)) ds \right)$$

An *Optimal Control* Perspective to Training

Loss Function

$$\ell(\theta, \mathbf{x}) := \underbrace{L(\mathbf{z}(S))}_{\text{Terminal Loss}} + \underbrace{\int_0^S l(s, \mathbf{z}(s)) ds}_{\text{Integral Loss}}$$

The training can be then cast into the *optimal control* problem

$$\begin{aligned} \min_{\theta \in \mathcal{W}} \quad & \mathbb{E}_{\mathbf{x}} [\ell(\theta, \mathbf{x})] \\ \text{subject to} \quad & \dot{\mathbf{z}}(s) = f_{\theta(s)}(s, \mathbf{x}, \mathbf{z}(s)) \quad s \in \mathcal{S} \\ & \mathbf{z}(0) = h_x(\mathbf{x}) \\ & \hat{\mathbf{y}}(s) = h_y(\mathbf{z}(s)) \end{aligned}$$

which we wish to solve via **stochastic gradient descent**.

Two options:

- backpropagate through the discrete steps of the solver
- backpropagate through the solution of the ODE

Can we compute gradients in closed-form?

Adjoint Method – Intuition

To retrieve the gradient, we have to “unroll” the integral
 \Rightarrow solve an other differential equation *backward* in \mathcal{S}

Generalized Adjoint: Finite-Dimensional Case

Let θ be *constant* in s [Chen T.Q. et al., 2018]

Proposition: Generalized Adjoint Gradients

Consider the loss function $\ell = L(\mathbf{z}(S)) + \int_S l(s, \mathbf{z}(s)) ds$. Then,

$$\frac{d\ell}{d\theta} = \int_S \mathbf{a}^\top(s) \frac{\partial f_\theta}{\partial \theta} ds \quad \text{where } \mathbf{a}(s) \text{ satisfies } \begin{cases} \dot{\mathbf{a}}^\top(s) = -\mathbf{a}^\top(s) \frac{\partial f_\theta}{\partial \mathbf{z}} - \frac{\partial l}{\partial \mathbf{z}} \\ \mathbf{a}^\top(S) = \frac{\partial L}{\partial \mathbf{z}(S)} \end{cases}$$

- This is the **exact** gradient;
- We do not need to store activations $\Rightarrow \mathcal{O}(1)$ memory efficiency

What about for a general class of $\theta(s)$?²

²Massaroli, Poli et al., Dissecting Neural ODEs, NeurIPS 2020

Why bother with the framework?

Advantages

- $\mathcal{O}(1)$ memory gradients
- Learning and Control
- Cheap Normalizing Flows
- Uncertainty Estimation
- Stability and Constraints

Disadvantages

- Expressivity limitations
- Compute requirements
- Novelty
- Different *tricks*
- Less accessible

Hypersolvers for Neural ODEs

Solving the Differential Equation

Solving differential equations is **costly**. Can be an issue for deployment.

Can we do better?

Yes. Analyze interplay between numerical solvers and Neural ODE.³
Ties in with pretraining strategies (NLP) and compression techniques.

Hot area at the moment. **Regularization techniques**⁴ also help by controlling the stiffness of learned ODEs, but they are not applicable in general.

³Poli, Massaroli et al., Hypersolvers: Toward Fast Continuous-Depth Models, NeurIPS 2020

⁴Finlay et al., How to train your neural ODE: the world of Jacobian and kinetic regularization, ICML 2020. Kelly et al., Learning Differential Equations that are Fast to Solve, NeurIPS2020

Quick Refresher on Explicit ODE Solvers

Iterate to solve:

$$\mathbf{z}_{k+1} = \mathbf{z}_k + \epsilon \psi(s_k, \mathbf{x}, \mathbf{z}_k)$$

- ODE solvers differ in how the map ψ is designed
- *Higher-order* (explicit) solvers compute ψ iteratively in p steps (p denotes the order of the solver)

[example] p -th order *Runge-Kutta* method:

$$\mathbf{r}_i = f_{\theta(s_k)}(s_k + \mathbf{c}_i \epsilon, \mathbf{x}, \mathbf{z}_k + \tilde{\mathbf{z}}_k^i) \quad i = 1, \dots, p$$

$$\tilde{\mathbf{z}}_k^i = \epsilon \sum_{j=1}^p \mathbf{a}_{ij} \mathbf{r}_j \quad i = 1, \dots, p$$

$$\psi = \sum_{j=1}^p \mathbf{b}_j \mathbf{r}_j$$

where $\mathbf{a} \in \mathbb{R}^{p \times p}$, $\mathbf{b} \in \mathbb{R}^p$, $\mathbf{c} \in \mathbb{R}^p$ fully characterize the method

Hypersolvers: main idea

Consider a p th order explicit ODE solver.

Discretized Neural ODE Solution

$$\begin{cases} \mathbf{z}_{k+1} = \mathbf{z}_k + \epsilon \psi(s_k, \mathbf{x}, \mathbf{z}_k) \\ \mathbf{z}_0 = h_x(\mathbf{x}) \\ \hat{\mathbf{y}}_k = h_y(\mathbf{z}_k) \end{cases} \quad k = 0, 1, \dots, K-1$$

ϵ : step size, ψ : solver's step.

Local Truncation Error

$$e_k := \|\mathbf{z}(s_{k+1}) - \mathbf{z}(s_k) - \epsilon \psi(s_k, \mathbf{x}, \mathbf{z}(s_k))\|_2 = \mathcal{O}(\epsilon^{p+1})$$

Hypersolver

$$\mathbf{z}_{k+1} = \mathbf{z}_k + \underbrace{\epsilon \psi(s_k, \mathbf{x}, \mathbf{z}_k)}_{\text{base solver step}} + \epsilon^{p+1} \overbrace{g_\omega(\epsilon, s_k, \mathbf{x}, \mathbf{z}_k)}^{\text{hypersolver net}}$$

Improvement on Solver Error

Given some *nominal* trajectories $\{(s_k, \mathbf{z}(s_k))\}_k$, consider the *residual*

$$\mathcal{R}_k = \frac{1}{\epsilon^{p+1}} [\mathbf{z}(s_{k+1}) - \mathbf{z}(s_k) - \epsilon \psi(\mathbf{x}, s_k, \mathbf{z}(s_k))]$$

Residual Training

Consider also the loss function

$$\ell = \frac{1}{K} \sum_{k=0}^{K-1} \|\mathcal{R}(s_k, \mathbf{z}(s_k), \mathbf{z}(s_{k+1})) - g_\omega(\epsilon, \mathbf{x}, s_k, \mathbf{z}(s_k))\|_2$$

We have that,

$$\forall k \quad \|\mathcal{R}_k - g_\omega\|_2 \leq \mathcal{O}(\delta) \quad \Rightarrow \quad e_k = \mathcal{O}(\delta \epsilon^{p+1})$$

At training, may use a lower-order solver to achieve the same accuracy!
 \Rightarrow Accelerate inference of Neural ODEs

Hypersolvers and Optimal Control

Hypersolvers for Optimal Control

Key points:

- We need accurate trajectories for training a controller
- If the optimization is done *online* it is even more crucial to speed up the simulation

⇒ Hypersolvers are helpful to obtain Pareto-optimal forward simulations!


We need to train the hypersolver on a reasonable set of the explored state-space and introduce the **control input** as an important feature

Training Hypersolvers for Controlled Systems

Residual Training for Controlled Systems

The control value u becomes an important input for the hypersolver

$$\ell = \frac{1}{K} \sum_{k=0}^{K-1} \left\| R(t_k, x(t_k), x(t_{k+1})) - g_{\omega}(t_k, x(t_k), \dot{x}(t_k), u(t_k)) \right\|_2$$



Direct Optimal Control

Direct Optimal Control

We want to obtain a control policy by *directly* optimizing the cost function over the complete trajectories

$$\begin{aligned} \min_{\mathbf{u}} \quad & J_{\mathbf{u}} \\ \text{subject to} \quad & \dot{x}(t) = f(t, x(t), \mathbf{u}(t)) \\ & x \in \mathbb{X}; \mathbf{u} \in \mathbb{U} \end{aligned}$$

where \mathbf{u} is the control policy and cost function J is

$$J_{\mathbf{u}} = x^{\top}(t_f) \mathbf{P} x(t_f) + \int_{t_0}^{t_f} [x^{\top}(t) \mathbf{Q} x(t) + \mathbf{u}^{\top}(t) \mathbf{R} \mathbf{u}(t)] dt$$

⇒ Hypersolvers make the simulation of whole trajectories more accurate and efficient compared to classical solvers!


Model Predictive Control

Model Predictive Control (MPC)

The following optimization problem is solved *online* at each time step:



$$\begin{aligned} \min_{\mathbf{u}} \quad & \sum_{i=1}^T J_{\mathbf{u}}(x(t_i)) \\ \text{subject to} \quad & \dot{x}(t) = f(t, x(t), \mathbf{u}(t)) \\ & x \in \mathbb{X}; \mathbf{u} \in \mathbb{U} \end{aligned}$$

\mathbf{u} : control policy; T : receding horizon in which future trajectories are predicted



Accelerating Model Predictive Control

MPC is an online optimization algorithm, it is even more crucial to obtain *fast and accurate* trajectory predictions



Results on a quadcopter model (17 states and 4 control variables)

Multi-stage Hypersolvers

If the dynamic model does not match perfectly the real one, we can use an **additional first-order term** to correct the vector field while the **second-order term** can be used to further improve accuracy

Multi-stage HyperEuler Step

$$x_{k+1} = x_k + \overbrace{\epsilon f(t_k, x_k, u_k)}^{\text{Euler step}} + \epsilon \overbrace{g'_w(t_k, x_k, u_k)}^{\text{1st order residual approximator}} + \epsilon^2 \overbrace{g''_w(t_k, x_k, u_k)}^{\text{2nd order residual approximator}}$$


Training Multiple Stages

If the *nominal* system model is unknown, we can not know what the residual directly: we can train directly on the known nominal trajectories

Training multi-stage hypersolvers


$$\ell = \frac{1}{K} \sum_{k=0}^{K-1} \text{MSE}(x(t_{k+1}), \hat{x}_{k+1})$$

where $x(t_{k+1})$ is obtained by the real (nominal) system and \hat{x}_{k+1} is one step of the Multi-stage HyperEuler



Extra: torchdyn

The new torchdyn 1.0 is out! Check it out at
<https://github.com/DiffEqML/torchdyn>



PyTorch library dedicated to neural differential equations and implicit models. Maintained by DiffEqML.

Thank you for your Attention!